

Structural and Semantic Synergy in 3D Scene Captioning: A Technical Blueprint for Million-Scale Grounded Understanding

The evolution of 3D vision-language (3D-VL) learning represents a critical frontier in the development of embodied agents capable of navigating and interacting with complex physical environments. While 2D vision-language models have achieved remarkable success through the ingestion of billion-scale datasets, the 3D domain has historically struggled with a profound scarcity of paired data and the lack of a unified framework for distilling grounded knowledge.¹ The task of 3D captioning, which requires an automated system to generate high-level natural language descriptions of objects and their spatial configurations within a 3D scan, sits at the intersection of geometric processing and linguistic reasoning.² To address the inherent complexities of this task, modern research has converged on two primary methodological pillars: graph-based approaches for modeling structural relationships and vocabulary-based approaches for unconstrained semantic alignment.³ This report provides an exhaustive technical analysis of these paradigms, with a particular focus on scaling strategies and architectural innovations.

Foundations of 3D Vision-Language Data Scaling

The primary bottleneck in 3D-VL has traditionally been the difficulty of collecting high-quality, aligned 3D-language data. Unlike 2D images, which can be scraped in massive quantities from the web, 3D data collection requires specialized scanning hardware and is subject to significant noise and occlusion.¹ Projects such as SCENEVERSE have sought to overcome this through the first million-scale 3D-VL dataset, curating 68,000 indoor scenes and 2.5 million vision-language pairs from both real-world sources and synthetic environments.¹ This represents an order-of-magnitude increase over previous benchmarks like ScanRefer and ReferIt3D.¹

Dataset	Scenes	Objects	Language Pairs	Annotation Method
ScanRefer ¹	1.5K	33K	52K	Human-annotated
ReferIt3D ¹	1.5K	33K	242K	Human-annotated
3D-LLM ¹	1.5K	186K	659K	Template/LLM Syn.
EmbodiedScan ¹	5K	890K	970K	Human/Syn.

SCENEVERSE ¹	68K	1.5M	2.5M	Multi-level Syn/Human
-------------------------	-----	------	------	--------------------------

The curation of SCENEVERSE involved unifying data from diverse sources, including real-world datasets like ScanNet, ARKitScenes, HM3D, and MultiScan, alongside synthetic scenes from Structured3D and ProcTHOR.¹ This synthesis required rigorous preprocessing, including room segmentation, point subsampling (to a maximum of 240,000 points per scene), axis alignment, and the mapping of all semantic labels to a standardized set of 607 classes.¹

Graph-Based Approaches for Structural Reasoning

Graph-based methods are fundamental to 3D scene captioning because they explicitly represent the underlying structure of the environment. In these models, a 3D scene is typically abstracted as a hierarchical graph $G = (V, E)$, where each node $v \in V$ represents a distinct 3D object instance—often parameterized by its centroid and bounding box—and edges $e \in E$ represent spatial or semantic relationships.³

Hierarchical Scene Graph Construction

The construction of 3D scene graphs follows a multi-stage pipeline designed to handle the noise inherent in point cloud data. The process begins with node instantiation based on instance annotations from point clouds.¹ Edges are determined by analyzing the vertical and horizontal proximity between these nodes. Research identifies 21 specific types of relationships categorized into four main groups: vertical in-contact (e.g., supported by, embedded into), vertical non-contact (e.g., hanging on, mounted on), horizontal proximity (e.g., near, in front of, to the left of), and multi-object relationships (e.g., between, aligned).¹

Category	Specific Relationships	Functional Implication
Vertical In-contact	Supported by, placed in, inside, embedded into	Foundational structural stability
Vertical Non-contact	Hanging on, mounted on, above, below, higher than	Affordance and mounting patterns
Horizontal	Left, right, in front of, behind, adjacent, next to	Spatial navigation and relative localization
Multi-object	Between, aligned	Complex grouping and layout symmetry

A significant innovation in graph-based modeling is the use of automated verification procedures. For example, if a model identifies a TV as "hanging on" a wall, an automated check based on common-sense mappings might refine this to "mounted on" to ensure linguistic naturalness.¹ Furthermore, hierarchical levels are determined by support relationships; objects on the floor constitute the base hierarchy, while objects supported by tables form the subsequent level.¹

Transitioning from GCNs to Transformers

Early iterations of graph-based 3D captioning relied on Graph Convolutional Networks (GCNs). However, GCNs are limited by the over-smoothing dilemma, where node features become indistinguishable as the number of layers increases, hindering the model's ability to capture global scene structure.³ To overcome this, contemporary models like SGFormer (Semantic Graph Transformer) employ Transformer layers as the building block.³

SGFormer introduces two specialized layers: the Graph Embedding Layer (GEL) and the Semantic Injection Layer (SIL).³ The GEL utilizes edge-aware self-attention to capture global representation among nodes while explicitly modeling inter-object relationships.³ This allows for global information passing, which is essential for understanding the complex intrinsic structural and semantic relationships within a 3D scene.³

Semantic Injection and LLM Enrichment

The Semantic Injection Layer (SIL) represents a convergence between graph modeling and large-scale language knowledge.³ Traditional models often rely solely on class names, which are linguistically sparse. SIL leverages Large Language Models (LLMs) like GPT-4 to expand category labels into detailed natural language descriptions (e.g., describing a bed's specific attributes and its typical relationship to surrounding items like nightstands).³ These enriched textual descriptions are encoded via a text encoder like CLIP and injected into the visual node features through a cross-attention mechanism.³

This approach is particularly effective in long-tail and zero-shot scenarios, where the visual features alone might be insufficient to differentiate rare object categories.³ By providing the model with "emergent" linguistic knowledge, it can distinguish between semantically similar but visually distinct objects, such as different types of industrial machinery in a workshop.³

Vocabulary-Based Approaches for Semantic Alignment

While graph-based methods focus on structure, vocabulary-based approaches aim for unconstrained semantic understanding. The core challenge in 3D-VL is aligning the geometric representation of a scene with a language space that can encompass an arbitrary number of concepts.¹⁰

Open-Vocabulary and Zero-Shot Learning

Open-vocabulary (OV) 3D scene understanding enables models to recognize objects and attributes beyond a fixed training taxonomy.⁹ This is typically achieved by aligning 3D visual features with the feature space of a pre-trained 2D vision-language model like CLIP.¹²

Frameworks such as OpenScene co-embed 3D point features with text strings and image pixels in the CLIP space.¹³ This alignment is facilitated by establishing associations between 3D points and pixels from posed multi-view images.¹²

Once aligned, the model can perform a variety of open-vocabulary tasks—such as 3D object search, material identification (e.g., "soft," "metal"), and room type classification—by simply calculating the cosine similarity between the 3D features and a user-provided text query.¹³

Vocabulary-Free and Training-Free Paradigms

A more recent development is Vocabulary-Free 3D Instance Segmentation (VoF3DIS), which removes the need for any predefined vocabulary or user query at test time.¹⁵ Models like PoVo prompt a vision-language assistant to retrieve a list of all objects present in multi-view images to form a per-scene vocabulary.¹⁵ This approach is inherently more flexible than traditional open-vocabulary methods, as it does not assume that the relevant concepts are already known to the user.¹⁵

Other frameworks, such as OpenVoxel, propose training-free algorithms for grouping and captioning sparse voxels.¹⁶ By leveraging the native capabilities of multi-modal large language models (MLLMs), OpenVoxel can produce meaningful group descriptions and caption every segment in a scene without the need for additional 3D supervision or the introduction of task-specific embeddings.¹⁶

Semantic Label Propagation and Mapping

A critical technical component of vocabulary-based approaches is the propagation of semantic labels across modalities. In JOPP-3D, RGB-D panoramic images are converted into 3D point clouds, and a 3D-to-panoramic semantic label propagation method is used via depth correspondence to enable multi-view-consistent semantic maps.¹¹ This ensures that the generated captions are consistent whether they are viewed from the perspective of a single camera or as part of the global 3D structure.¹¹

Paradigm	Key Mechanism	Training Requirement	Flexibility
Closed-Vocabulary	Fixed label prediction (e.g., SoftGroup)	Full 3D supervision	Limited
Open-Vocabulary	CLIP-aligned 3D-text embedding (OpenScene)	2D-3D distillation	High
Vocabulary-Free	LLM-driven scene discovery (PoVo)	Training-free/inference-driven	Maximal
Training-Free	Direct MLLM captioning (OpenVoxel)	None	High

Architectural Synergy: Integrating Graph and Vocabulary Methods

The most effective 3D captioning models do not rely on a single methodology but rather synthesize graph-based structural reasoning with vocabulary-based semantic alignment. This synergy is best exemplified in the Grounded Pre-training for Scenes (GPS) framework and the

3DGraphLLM method.¹

Multi-Level Contrastive Alignment (GPS)

The GPS model, trained on the SCENEVERSE dataset, employs multi-level contrastive losses to align 3D scenes and text at various granularities.¹ The architecture decomposes the problem into three distinct levels of grounding:

1. **Object-Level Grounding:** This stage utilizes an object point cloud encoder (often PointNet++) to extract features, which are then aligned with object-level captions through a contrastive objective similar to CLIP.¹ For instances where captions are missing, class labels are used to construct simple templates like "a point cloud of".¹
2. **Referral-Object-Level Grounding:** To model the relationships crucial for referring expressions, GPS uses a self-attention-based reasoning transformer.¹ This module takes scene-object features and an object referral text, performing intra-scene contrast to force the selection of positive object-text pairs within the same environment.¹ This mimics the fine-grained region-word alignment successful in 2D-VL models.¹
3. **Scene-Level Grounding:** A spatial-attention transformer aggregates object features by incorporating their spatial locations.¹ The resulting scene feature is max-pooled and aligned with a global scene caption.¹

Learnable Scene Graph Representations for LLMs

3DGraphLLM takes the synergy a step further by constructing a learnable representation of a 3D scene graph specifically as an input to Large Language Models.⁵ Instead of representing a scene as a flat list of objects or a complex geometric point cloud, 3DGraphLLM encodes the scene as a structured sequence of subgraphs.⁵

The method employs an algorithm to select object subgraphs using k-nearest neighbors (k-NN) with Non-Maximum Suppression (NMS) to reduce redundancy.⁵ These subgraphs are provided to the LLM as a sequence of triplets, where semantic relations are encoded using features derived from the semantic edges of the scene graph.⁵ Research shows that explicitly using semantic relationships—such as those generated by state-of-the-art methods like VL-SAT—significantly improves the quality of LLM responses in tasks like dense scene captioning and question answering.⁵

Architecture Type	Relational Module	Language Interface	Key Task Advantage
Detect-then-describe (Scan2Cap)	Graph-based module	GRU/LSTM decoder	Simple sequential logic ²¹
Set-to-set (Vote2Cap-DETR)	Self-attention queries	Transformer decoder	Parallelism, reduced error propagation ²³
Graph-to-LLM (3DGraphLLM)	Semantic subgraphs	LLM prompts (triplets)	Complex reasoning, faster inference ⁵
Multi-level GPS	Spatial Attention	Contrastive Alignment	Zero-shot transfer, data

	Transformer		efficiency ¹
--	-------------	--	-------------------------

Automated Generation and LLM-Refined Language

A cornerstone of modern 3D captioning is the reliance on automated generation pipelines to create the vast amounts of text required for pre-training. SCENEVERSE utilizes a novel pipeline that leverages 3D scene graphs and LLMs to create multi-level descriptions.¹

Multiview-Aware Object Captioning

To generate detailed descriptions of an object's visual properties, the system identifies the object's occurrence in multi-view images through rendering.¹ The images are cropped based on rendered bounding boxes and processed through a model like BLIP2.¹ To ensure coherence, the system selects the top 10 initial captions with the highest CLIP similarity score and minimal occlusion.¹ These are then fed into an LLM to generate a summarized, error-corrected object caption (e.g., identifying a nightstand by its white color and its typical association with a laptop or bag).¹

Relationship and Scene Captioning

Object relationship captions (referrals) are generated by extracting spatial relationship triplets from the constructed 3D scene graph.¹ Various templates spanning passive and active tenses—as well as inversion clauses—are used to create initial descriptions.¹ An LLM is then employed to rephrase these into more natural, context-rich sentences.¹ For example, a template-based sentence "the chair is next to the sofa" might be rephrased to "the chair is situated beside the plush couch, creating a cozy seating area".¹ For scene-level captions, the LLM is provided with the room type, object counts, and a random subset of nodes and edges from the scene graph.¹ This ensures diversity in the generated global descriptions, focusing on different functional aspects of the room each time.¹

Performance Analysis and Empirical Benchmarking

The efficacy of scaling 3D vision-language learning is demonstrated through rigorous evaluation on established benchmarks. The data-scaling effect is consistent across tasks: as the volume of training data increases, both zero-shot and fine-tuned performance improve significantly.¹

Comparison of Visual Grounding Accuracies

GPS achieves state-of-the-art results on three major benchmarks: ScanRefer, Nr3D, and Sr3D.¹ Notably, even without additional fine-tuning, the pre-trained GPS model outperforms existing baselines on the ScanRefer dataset.¹

Method	Nr3D Overall	Sr3D Overall	ScanRefer Acc@0.5
3DVG-Trans ¹	40.8	51.4	34.7
SAT ¹	56.5	57.9	50.8

VIL3DRel ¹	64.4	72.8	68.6
3D-VisTA ¹	64.2	76.4	75.1
Ours (GPS scratch)	58.7	68.4	71.3
Ours (GPS pre-train)	55.2	74.1	77.4
Ours (GPS fine-tuned)	64.9	77.5	77.9

The results suggest that while complex structural designs (Ours scratch) are useful, they are no substitute for the scaling effect provided by million-scale data (Ours pre-train).¹ The transition from scratch training to pre-training on SCENEVERSE provides a substantial boost, particularly in the Accuracy@0.5 metric, indicating improved spatial localization capabilities.¹

Zero-Shot Generalization and Domain Transfer

One of the most compelling pieces of evidence for the success of these methods is their zero-shot transfer capability. Models trained on SCENEVERSE (excluding target dataset scenes) demonstrate superior generalization to unseen environments compared to models trained on single datasets.¹

The zero-shot performance is further enhanced in the "zero-shot text" setting, where models are trained on all scenes but with exclusively unseen scene-text distributions.¹ This indicates that the automatically generated scene-text pairs in SCENEVERSE supply sufficient knowledge for the model to comprehend a wide variety of scene distributions and linguistic styles.¹

Multi-Faceted Understanding: QA and Segmentation

The scaling effect extends beyond grounding and captioning to question answering (QA) and semantic segmentation.¹ GPS achieves state-of-the-art results on ScanQA and SQA3D, outperforming strong baselines like 3D-LLM.¹ In open-vocabulary semantic segmentation, pre-training the RegionPLC model on SCENEVERSE leads to consistent improvements in mIoU and mAcc on the ScanNet dataset.¹

Furthermore, for closed-vocabulary semantic segmentation, training the SWIN3D model on SCENEVERSE results in a 6% performance improvement over the original implementation.¹ This serves as strong evidence that the diverse scene data and multi-level language descriptions in SCENEVERSE capture fundamental geometric and semantic patterns applicable to a wide range of 3D understanding tasks.¹

Technical Insights for Future Model Design

Analyzing the success of million-scale 3D-VL models yields several key insights for future architectural development.

The Importance of Intra-Scene Contrast

The referral-object-level grounding objective in GPS highlights the critical role of intra-scene contrast.¹ By forcing the model to select the correct object from among multiple candidates in the same environment, it develops a much finer understanding of spatial relationships.¹ This is a

departure from earlier models that focused primarily on inter-scene contrast (distinguishing between different rooms) and is a primary driver of the model's success in complex referring expression tasks.¹

Addressing the Domain Gap: Real vs. Synthetic

Experimental analysis reveals a persistent domain gap between real and synthetic data.¹ While synthetic environments (e.g., Structured3D, ProcTHOR) offer unlimited scalability, models trained purely on these data suffer significant performance drops when transferred to real scenes like ScanNet.¹ Conversely, models trained on real scenes exhibit much better generalization to synthetic environments.¹ This suggests that the rate-limiting step for the next generation of 3D-VL scaling is not just the number of scenes, but the collection of high-quality, realistic scenes that capture natural 3D distributions.¹

Functional and Interaction-Aware Modeling

The shift toward functional 3D scene graphs (e.g., OpenFunGraph) indicates that the future of 3D captioning is not just descriptive but actionable.⁸ By modeling the functional relationships between objects and their interactive parts (like handles, switches, and remotes), models can provide information that is immediately useful for robotic navigation and manipulation.⁸ For instance, a system can correctly infer that a light switch controls a specific ceiling light rather than a nearby table lamp, enabling complex task-oriented interactions in indoor spaces.²⁷

Efficiency and Token Compression

As MLLMs become central to 3D scene understanding, inference efficiency is becoming a major bottleneck.²⁹ Techniques such as the learnable graph embeddings in 3DGraphLLM—which use five times fewer tokens than textual descriptions—are essential for real-time applications.⁵ The use of transformer-based compressors that receive aligned embeddings and output retention probabilities represents a promising direction for handling high-resolution 3D data efficiently.²⁹

Conclusion

The advancement of 3D scene captioning relies on a dual strategy of structural graph modeling and large-scale semantic alignment. The integration of graph-based approaches—such as the Hierarchical Open-Vocabulary Scene Graphs (HOV-SG) and the Semantic Graph Transformer (SGFormer)—allows for a deep understanding of the complex, multi-level relationships that define a scene's layout.³ Simultaneously, vocabulary-based methods provide the unconstrained semantic breadth necessary to describe an open-world environment.⁹

The transition to million-scale datasets like SCENEVERSE, coupled with unified frameworks like GPS, has set a new paradigm for 3D vision-language grounding.¹ These models demonstrate that large-scale pre-training on multi-level scene-language pairs can result in emerging zero-shot capabilities and state-of-the-art performance across a diverse set of downstream tasks, from grounding to question answering and semantic segmentation.¹ Future developments are poised to further refine these capabilities by focusing on fine-grained part-level details,

industrial domain adaptation, and functionally actionable scene representations, ultimately paving the way for intelligent agents that can understand and interact with the 3D world with human-like proficiency.⁹

Works cited

1. btp_paper.pdf
2. 3D Object Captioning - Emergent Mind, accessed on April 27, 2026, <https://www.emergentmind.com/topics/3d-object-captioning>
3. SGFormer: Semantic Graph Transformer for Point Cloud-Based 3D Scene Graph Generation, accessed on April 27, 2026, <https://ojs.aaai.org/index.php/AAAI/article/view/28197/28391>
4. Open3DSG: Open-Vocabulary 3D Scene Graphs from Point Clouds with Queryable Objects and Open-Set Relationships - arXiv, accessed on April 27, 2026, <https://arxiv.org/html/2402.12259v2>
5. 3DGraphLLM: Combining Semantic Graphs and Large Language Models for 3D Scene Understanding - CVF Open Access, accessed on April 27, 2026, https://openaccess.thecvf.com/content/ICCV2025/papers/Zemskova_3DGraphLLM_Combining_Semantic_Graphs_and_Large_Language_Models_for_3D_ICCV_2025_paper.pdf
6. 3D Semantic Label Results - ScanNet Benchmark, accessed on April 27, 2026, https://kaldir.vc.in.tum.de/scannet_benchmark/semantic_label_3d
7. arxiv.org, accessed on April 27, 2026, <https://arxiv.org/html/2412.18450v3>
8. Functional 3D Scene Graphs in AI & Robotics - Emergent Mind, accessed on April 27, 2026, <https://www.emergentmind.com/topics/functional-3d-scene-graphs>
9. Open-vocabulary 3D scene perception in industrial environments - arXiv, accessed on April 27, 2026, <https://arxiv.org/html/2602.19823v1>
10. Hierarchical Open-Vocabulary 3D Scene Graphs for Language-Grounded Robot Navigation, accessed on April 27, 2026, <https://arxiv.org/html/2403.17846v2>
11. JOPP-3D: Joint Open Vocabulary Semantic Segmentation on Point Clouds and Panoramas, accessed on April 27, 2026, <https://arxiv.org/html/2603.06168v2>
12. Group3D: MLLM-Driven Semantic Grouping for Open-Vocabulary 3D Object Detection, accessed on April 27, 2026, <https://arxiv.org/html/2603.21944v1>
13. OpenScene: 3D Scene Understanding with Open Vocabularies, accessed on April 27, 2026, <https://arxiv.org/pdf/2211.15654>
14. CVPR Poster OpenScene: 3D Scene Understanding With Open Vocabularies, accessed on April 27, 2026, <https://cvpr.thecvf.com/virtual/2023/poster/22442>
15. Vocabulary-Free 3D Instance Segmentation with Vision and Language Assistant - arXiv, accessed on April 27, 2026, <https://arxiv.org/html/2408.10652v1>
16. [2601.09575] OpenVoxel: Training-Free Grouping and Captioning Voxels for Open-Vocabulary 3D Scene Understanding - arXiv, accessed on April 27, 2026, <https://arxiv.org/abs/2601.09575>
17. OpenVoxel: Training-Free Grouping and Captioning Voxels for Open-Vocabulary 3D Scene Understanding | NVIDIA Research Taiwan, accessed on April 27, 2026, https://research.nvidia.com/labs/twn/publication/cvpr_2026_openvoxel/

18. OpenVoxel: Training-Free Grouping and Captioning Voxels for Open-Vocabulary 3D Scene Understanding - arXiv, accessed on April 27, 2026, <https://arxiv.org/html/2601.09575v1>
19. OpenVoxel: Training-Free Grouping and Captioning Voxels for Open-Vocabulary 3D Scene Understanding - Hugging Face, accessed on April 27, 2026, <https://huggingface.co/papers/2601.09575>
20. JOPP-3D: Joint Open Vocabulary Semantic Segmentation on Point Clouds and Panoramas, accessed on April 27, 2026, https://www.researchgate.net/publication/401692417_JOPP-3D_Joint_Open_Vocabulary_Semantic_Segmentation_on_Point_Clouds_and_Panoramas
21. TOD3Cap: Towards 3D Dense Captioning in Outdoor Scenes - ECVA | European Computer Vision Association, accessed on April 27, 2026, https://www.ecva.net/papers/eccv_2024/papers_ECCV/papers/02800.pdf
22. Bi-directional Contextual Attention for 3D Dense Captioning - ECVA, accessed on April 27, 2026, https://www.ecva.net/papers/eccv_2024/papers_ECCV/papers/02804.pdf
23. End-to-End 3D Dense Captioning With Vote2Cap-DETR - A*OAR, accessed on April 27, 2026, <https://oar.a-star.edu.sg/storage/x/xom0w6jv2x/chen-end-to-end-3d-dense-captioning-with-vote2cap-detr-cvpr-2023-paper.pdf>
24. CVPR Poster End-to-End 3D Dense Captioning With Vote2Cap-DETR, accessed on April 27, 2026, <https://cvpr.thecvf.com/virtual/2023/poster/22688>
25. 3DGraphLLM: Combining Semantic Graphs and Large Language Models for 3D Scene Understanding. - YouTube, accessed on April 27, 2026, <https://www.youtube.com/watch?v=m7gpP0oZhPc>
26. 3D-VisTA: Pre-trained Transformer for 3D Vision and Text Alignment - CVF Open Access, accessed on April 27, 2026, https://openaccess.thecvf.com/content/ICCV2023/papers/Zhu_3D-VisTA_Pre-trained_Transformer_for_3D_Vision_and_Text_Alignment_ICCV_2023_paper.pdf
27. Open-Vocabulary Functional 3D Scene Graphs for Real-World Indoor Spaces - Liner, accessed on April 27, 2026, <https://liner.com/review/openvocabulary-functional-3d-scene-graphs-for-realworld-indoor-spaces>
28. CVPR Poster Open-Vocabulary Functional 3D Scene Graphs for Real-World Indoor Spaces, accessed on April 27, 2026, <https://cvpr.thecvf.com/virtual/2025/poster/34720>
29. EvoComp: Learning Visual Token Compression for Multimodal Large Language Models via Semantic-Guided Evolutionary Labeling - arXiv, accessed on April 27, 2026, <https://arxiv.org/html/2604.17087v1>
30. 3D-VisTA: Pre-trained Transformer for 3D Vision and Text Alignment [Quick Review] - Liner, accessed on April 27, 2026, <https://liner.com/review/3dvista-pretrained-transformer-for-3d-vision-and-text-alignment>
31. ExCap3D: Expressive 3D Scene Understanding via Object Captioning with Varying Detail - CVF Open Access, accessed on April 27, 2026,

https://openaccess.thecvf.com/content/ICCV2025/papers/Yeshwanth_ExCap3D_Expressive_3D_Scene_Understanding_via_Object_Captioning_with_Varying_ICCV_2025_paper.pdf